

---

THE GEORGE  
WASHINGTON  
UNIVERSITY  

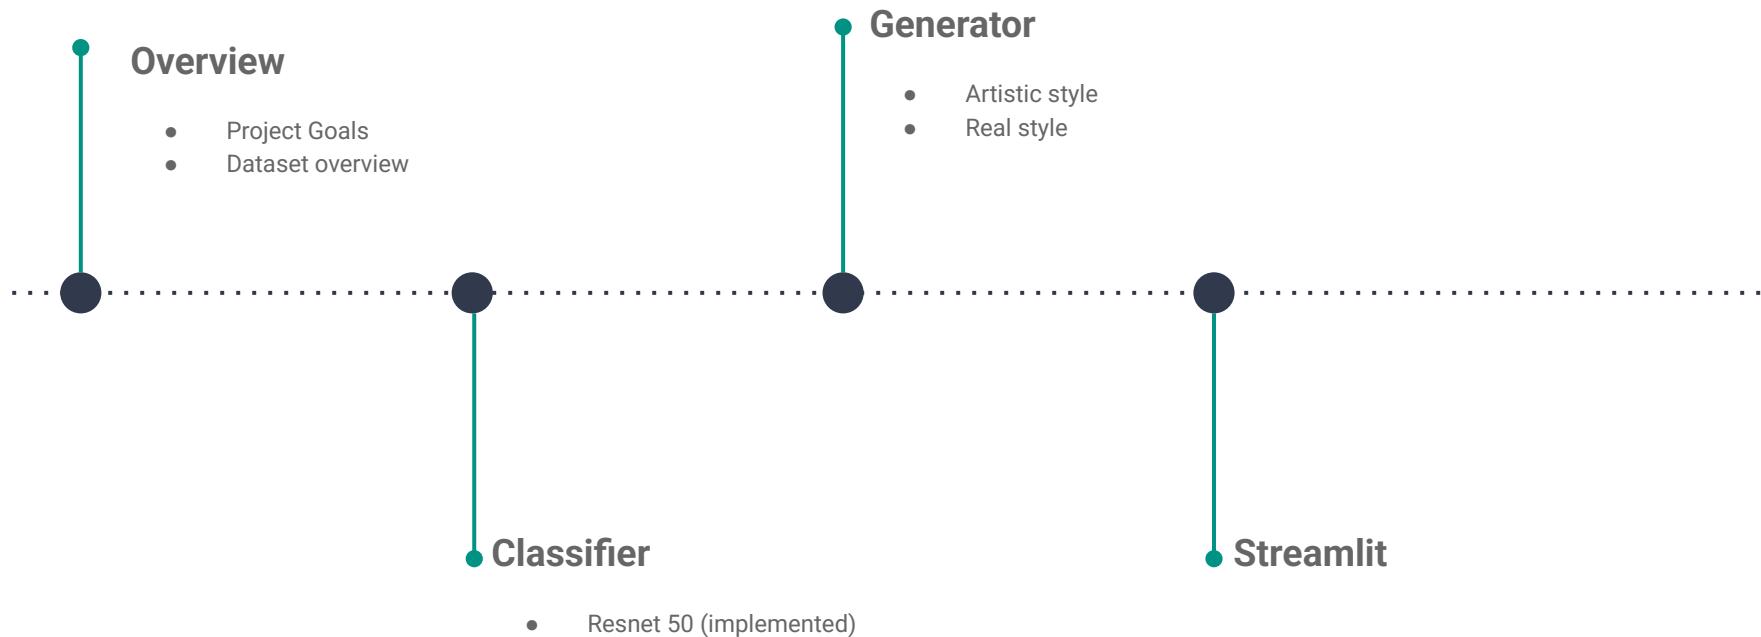
---

WASHINGTON, DC

# Dog Breed Classification and Generation

Adam Stuhltrager | Juhua Deng | Sameer Batra | Sayan Patra

# Table of Contents



# Project Goals

- **Develop a Highly Accurate Dog-Breed Classifier**

Train and evaluate a ResNet-50 model capable of reliably distinguishing between **120 dog breeds**, even those with subtle visual differences.

- **Build a Photorealistic Dog Image Generator**

Create an SDXL + LoRA-powered generator that can synthesize **anatomically consistent, breed-faithful** dog images in multiple visual styles.

- **Unify Classification and Generation into a Single System**

Integrate the classifier and generator so the predicted breed can automatically guide or enhance image synthesis, forming a cohesive recognition-to-generation pipeline.

- **Ensure Strong Quantitative and Qualitative Performance**

Optimize both components to achieve **high validation accuracy**, strong generalization, and visually realistic generation output across diverse breeds and poses.

- **Deploy an Interactive Application for Real-World Use**

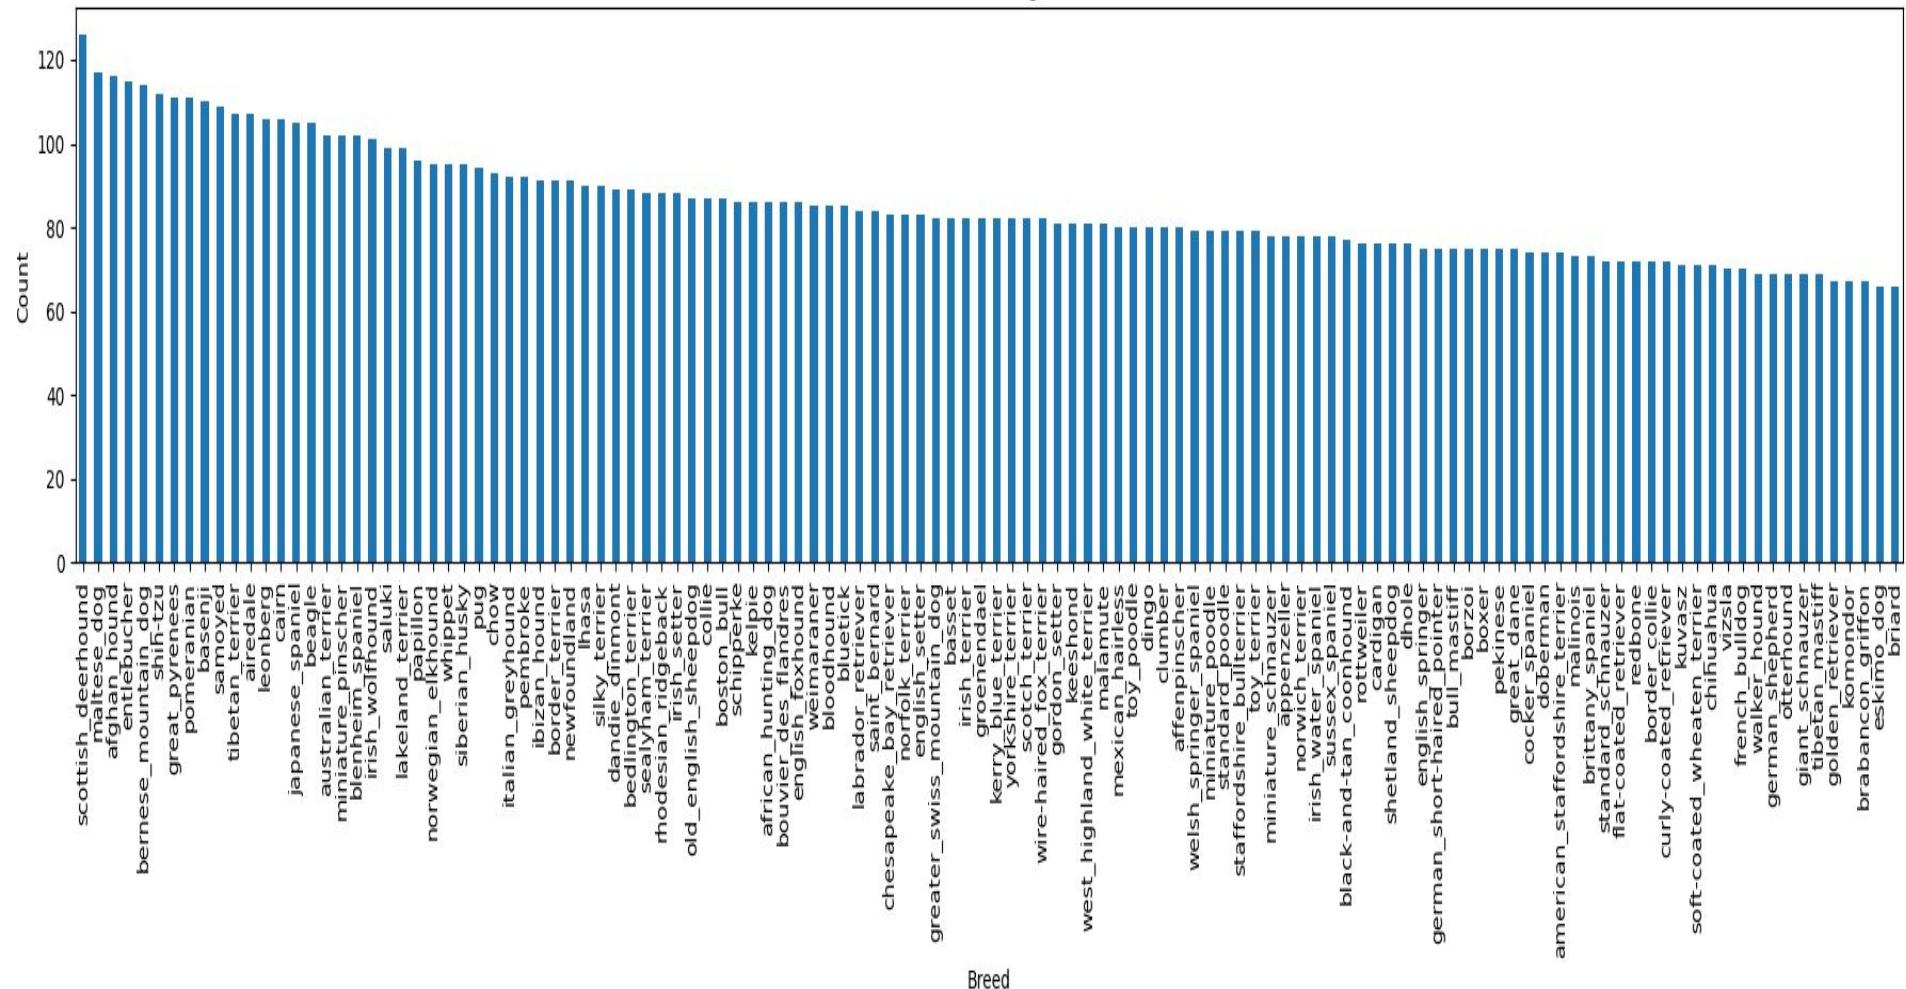
Build a Streamlit-based interface that allows users to **upload images, identify breeds, and generate custom dog images** in real time, demonstrating full system functionality.

# Dataset Overview

## Dog Breed Identification Dataset

- 10,222 Images
- Number of classes: 120 dog breeds
- Minimum Images per breed: ~60
- Maximum Images per breed: ~130

# Distribution of Dog Breeds



Classifier

# Classifier

## Data Preprocessing

- **Unified Image Format:**  
All inputs are converted to **RGB** to ensure consistent 3-channel structure, avoiding mismatches with the pretrained ResNet-50 expectations.
- **Spatial Normalization:**  
Images are resized to **224×224** followed by a **center crop** to remove spatial distortions and maintain uniform input size for convolution layers.
- **Tensor Conversion:**  
Images are transformed into **PyTorch tensors**, enabling batched GPU computation, gradient flow, and compatibility with the model pipeline.
- **Statistical Normalization:**  
Pixel values are standardized using **ImageNet mean/std** ( $\mu=[0.485, 0.456, 0.406]$ ,  $\sigma=[0.229, 0.224, 0.225]$ ) to align with pre trained feature distributions.
- **Transformation Consistency:**  
The exact validation preprocessing pipeline is replicated during inference to ensure **no distribution shift** between training and prediction.

# Classifier

## Data Loader Configuration

- **Batch Size Optimization:**

Batch sizes of **128** balance memory efficiency and gradient stability, improving GPU utilization during training.

- **Randomized Training Batches:**

Setting “`shuffle=True`” for training ensures each epoch sees a different sample order, enhancing generalization and preventing memorization.

- **Deterministic Evaluation:**

Validation/testing Data Loaders keep “`shuffle=False`” to ensure **consistent, repeatable** evaluation metrics.

- **Parallel Data Loading:**

Using **4–8 workers** loads multiple images simultaneously, reducing I/O bottlenecks and keeping GPUs fully utilized.

- **Pinned Memory Acceleration:**

“`pin_memory=True`” speeds up CPU to GPU transfers by locking memory pages, enabling high-throughput batch delivery.

# Classifier

## Learning Rate Scheduler

### **Dynamic LR Adjustment:**

Schedulers modify the learning rate during training to allow **fast initial learning** and **stable fine-tuning** as gradients get smaller.

### **Avoiding Optimization Pitfalls:**

LR decay prevents the optimizer from overshooting minima or oscillating, especially important for deep networks like ResNet-50.

### **Scheduler Options:**

- **Step LR:** periodic LR drops
- **CosineAnnealing:** smooth cyclic decay
- **ReduceLROnPlateau:** triggered when validation loss stagnates

### **Better Convergence Quality:**

Lower LR in later epochs enables the model to fine-tune high-level features, improving classification accuracy on similar dog breeds.

### **Training Stability:**

Controlled LR decay reduces gradient explosion risk and ensures smoother loss curves throughout training.

# Classifier

## Early Stopping

- **Continuous Validation Monitoring:**  
The system tracks **validation loss each epoch**, acting as a signal of generalization performance.
- **Patience-Based Stopping:**  
Training halts when the model fails to improve for a specified number of epochs, preventing wasteful computation.
- **Overfitting Prevention:**  
Early stopping stops learning before training noise is memorized, enhancing real-world predictive robustness.
- **Best Checkpoint Preservation:**  
The checkpoint with the lowest validation loss is retained, ensuring deployment uses the **optimal** model version.
- **Efficient Training Time:**  
Cuts off unnecessary training epochs, especially when improvements plateau, enabling faster experimentation cycles.

# Classifier

## Integrated Training Workflow

- **Input Standardization:**  
Preprocessing ensures images match the input distribution expected by pretrained ResNet-50, improving feature extraction fidelity.
- **Data Feeding Efficiency:**  
An optimized Data Loader pipeline ensures smooth, non-blocking GPU usage, maintaining high training throughput.
- **Adaptive Learning Control:**  
Learning rate schedulers dynamically adjust optimization behavior, stabilizing convergence and improving accuracy.
- **Generalization Safeguards:**  
Early stopping protects against overfitting and ensures the best-performing model is always selected.
- **Robust Final Model:**  
Together, these components form a consistent, efficient, and high-performing training loop for dog-breed classification.

# Classifier

## Model Architecture

- **Backbone: ResNet-50 Deep CNN**  
Utilizes a 50-layer Residual Network consisting of convolutional layers with skip connections, enabling stable training of deep architectures by mitigating vanishing gradients.
- **Bottleneck Residual Blocks**  
Each block uses a  $1 \times 1 \rightarrow 3 \times 3 \rightarrow 1 \times 1$  convolution structure, reducing computational cost while allowing extraction of fine-grained texture patterns crucial for distinguishing similar dog breeds.
- **Global Average Pooling Layer**  
Compresses spatial feature maps into a single vector per channel, making the network robust to object location and improving breed classification accuracy.
- **Custom Fully Connected Classifier Head**  
The original ImageNet FC layer is replaced with a **Linear(`num_features` → `num_classes`)** layer aligned with the exact number of dog breeds, enabling fine-tuned classification.
- **Softmax Output & Confidence Scoring**  
Final logits are passed through a softmax layer to produce breed-wise probabilities, allowing extraction of both **top-1 prediction** and **confidence score** for deployment and UI integration.

# Classifier

## Results

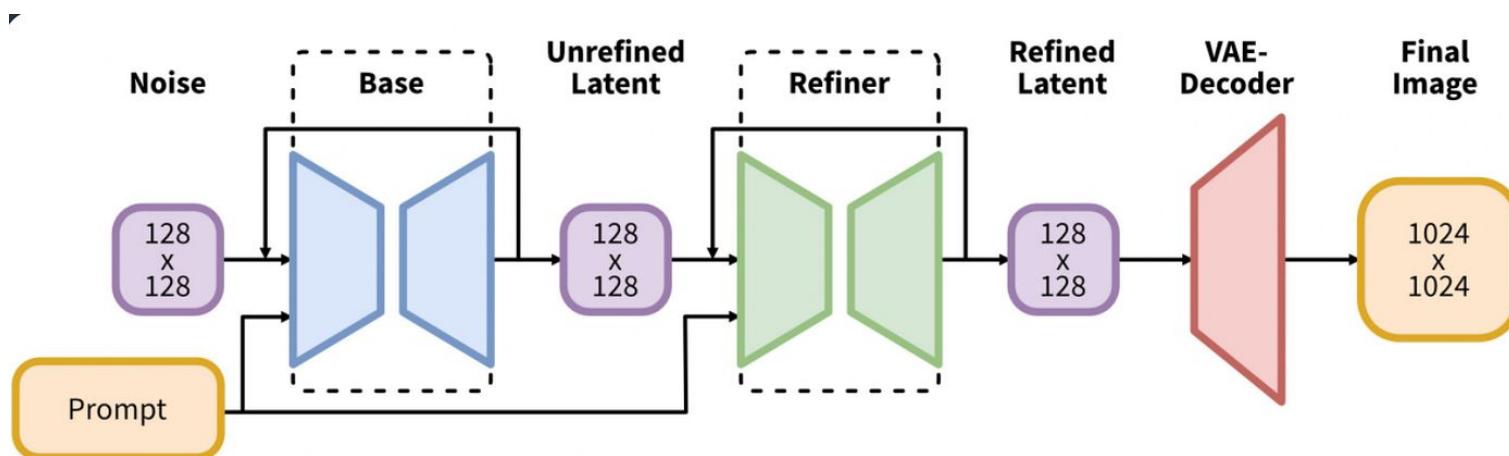
	Macro	Micro	Weighted
Precision	0.8508	0.8508	0.8601
Recall	0.8461	0.8508	0.8508
F1-score	0.8435	0.8508	0.8508
Accuracy			0.8508

Generator

# Generator

## SDXL

Using a large dual-UNet diffusion architecture that gradually denoises latent representations into high-resolution images. It combines two powerful text encoders to understand prompts more precisely, and its two-stage generation process (Base + Refiner) allows the model to capture global structure first and then add fine-grained details.



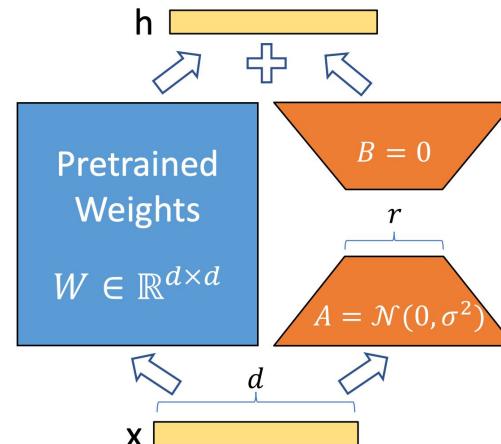
# Generator

	SDXL	SD 1.5	SD 2.1
Original Resolution	1024x1024	512x512	768x768
Parameters	2.6B	0.86B	0.86B
Visual Quality	Significantly higher fidelity	Moderate	Improved over 1.5
Prompt Adherence	Excellent	Moderate	Good

# Generator

## LoRA

An efficient fine-tuning method that modifies a large model by adding a small low-rank update to its original weight matrices. Instead of training the full weight  $W$ , LoRA freezes  $W$  and learns an update represented as



# Generator

## Model Loading & Initialization

- **Base Model Download:**  
Loads SDXL via `StableDiffusionXLPipeline.from_pretrained()` using either fp16 (GPU) or fp32 (CPU).
- **torch\_dtype Selection:**  
Dynamically sets dtype to `float16` on CUDA for speed and memory efficiency; float32 on CPU for compatibility.
- **LoRA Repository Handling:**  
LoRA repo & weight names are configurable:  
`lora_repo`, `lora_weight_name`, `lora_scale`.
- **Style Customization:**  
LoRA scaling enables mixing SDXL's base model with stylistic layers (e.g., manga, cinematic, line-art).

# Generator

## Prompt Engineering

- **Primary Prompt (Limit to 77 tokens):**  
Accepts breed-specific text and adds detailed descriptors such as **anatomy**, **pose**, **style**, **linework**, etc.
- **Negative Prompt:**  
Carefully crafted to prevent common SDXL artifacts:  
extra tails, extra limbs, distorted anatomy, CGI look, blurred areas, missing legs.
- **Dog Anatomy Optimization:**  
Includes constraints like “*single visible tail*”, “*proper proportions*”, “*natural limb spacing*” to avoid morphological hallucinations.
- **Style Control:**  
Can be switched between Manga/Anime/Pastel Anime/Pixel Art.

# Generator

Negative Examples:



Multiple Tails



Extra Limbs

# Generator

## Settings

- **Seeded Generator:**  
Uses `torch.Generator` to enforce deterministic output when a `seed` is provided.
- **Additional Cross-Attention Parameters:**  
If LoRA is active, the generator injects `"cross_attention_kwargs": {"scale": lora_scale}` to blend styles smoothly.
- **Sampling Configuration:**
  - Resolution: **1080 × 720**
  - Inference Steps: **50**
  - Guidance Scale: **5.0**  
These balance quality, consistency, and style strength.
- **Pipeline Call:**  
Executes `self.pipe()` with prompt, negative prompt, shapes, guidance, generator, and LoRA scaling.

# Generator

## Realistic SDXL Generation Overview

- **Text-to-Latent Pipeline Creation**

The system constructs a two-stage SDXL pipeline (Base → Refiner), enabling structured latent formation first, followed by refinement, mimicking how real image formation happens incrementally.

- **Anatomy-Grounded Prompt Injection**

Prompts are dynamically built from breed-specific anatomical descriptors and global biomechanical priors, ensuring that SDXL starts with correctly conditioned structural information before sampling begins.

- **Negative Prompt Enforcement**

Combined negative prompts strictly penalize SDXL's common anatomical failure modes (extra limbs, fused paws, warped joints), forcing the pipeline to remain anatomically aligned during diffusion steps.

- **Sequential Denoising (Base → Refiner) Workflow**

The Base model generates the global structure in latent space, while the Refiner model sharpens texture, fur detail, paw edges, and lighting, functioning as a second “micro-detail” denoising phase.

- **Real-Life Post-Processing Integration**

After SDXL generates the refined image, custom physical post-processing (microfur, tone curve, sensor noise) is applied to align the final output with real camera characteristics.

# Generator

## Prompt Engineering Pipeline

- **Breed-Specific Anatomy Extraction**  
The system retrieves skull, fur, paws, tail, and body descriptions from BREED\_ANATOMY and merges them with hind-leg biomechanics for structurally correct prompt conditioning.
- **Generic Prior Fallback Mechanism**  
When the breed is not found, the pipeline replaces missing details with a robust global canine anatomy prior, ensuring SDXL always receives sufficiently detailed geometric constraints.
- **Photographic Prior Injection (PHOTO\_PRIOR)**  
Every prompt is reinforced with DSLR-specific cues—lens focal length, RAW texture behavior, dynamic range—which trains SDXL to hallucinate a *photograph*, not a stylized render.
- **Pose Conditioning and Leg Visibility**  
Full-body side-view constraints ensure that SDXL must render all four limbs, distribute weight naturally, and expose the hind-leg structure — a common failure area corrected here.
- **Unified Prompt Chain for Base and Refiner**  
Both SDXL stages share the same positive/negative prompt embeddings, ensuring consistent structural intent across denoising phases.

# Generator

## Sampling & Latent Pipeline

- **Timesteps and Scheduler Configuration**

The Euler Ancestral scheduler creates a stable noise-to-latent trajectory, balancing sharp detail with anatomical accuracy across the Base and Refiner phases.

- **Denoising Split Strategy**

Total steps are divided between the Base (structure) and Refiner (detail). The Base denoises until `denoising_end`, handing the latents to the Refiner at the correct “noise depth.”

- **Prompt Embedding Encoding**

Both positive and negative prompts are encoded into high-dimensional embeddings. SDXL uses these embeddings as conditioning at every timestep through cross-attention.

- **Latent Representation Stabilization**

Intermediate latent images are generated in a noise-controlled latent space, enabling SDXL to maintain consistent anatomy before translating them into pixel-space.

- **Refiner Enhancement Pipeline**

The refiner continues denoising only the last portion of the pipeline, sharpening hair strands, eyes, paw textures, and structural edges without altering global anatomy.

# Generator

## Anatomy-Control Pipeline

- **Hind-Leg Biomechanics Enforcement**

The generator inserts detailed biomechanical descriptors for femur–tibia proportions, stifle angles, and hock behavior to force SDXL into anatomically valid limb representation.

- **Breed-Specific Additive Anatomy**

Skull shape, paw types, coat density, tail carriage, ear posture, and gait descriptions are appended to the prompt, giving SDXL explicit structural targets.

- **Global Canine Anatomy Stack**

Even if breed-specific details fail, the pipeline always includes the general canine skeletal and muscular structure, preventing SDXL from drifting into iconographic or stylized forms.

- **Hard Negative Penalties for Limb Failures**

Negative prompts explicitly describe unacceptable outputs (fused paws, extra legs, misplaced joints), shaping the loss landscape to punish structural hallucinations.

- **Prompt Embedding Uniformity Across SDXL Stages**

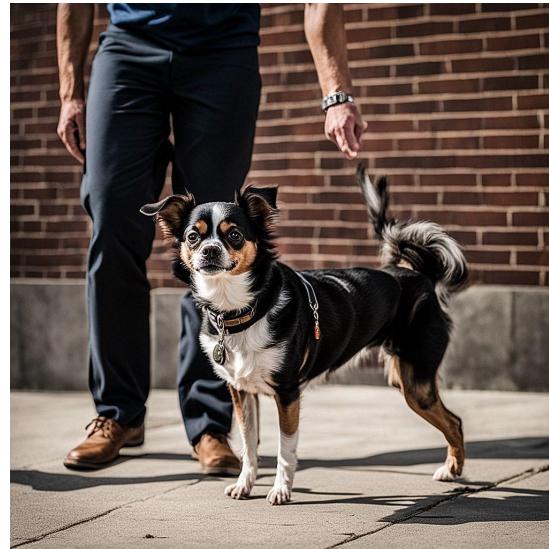
Ensures anatomy constraints affect both the early low-frequency structure (Base) and final micro-detail refinement (Refiner), maintaining consistency throughout the pipeline.

# Generator

## Negative Examples



Incorrect perspective, multiple limbs



Lack of limb



Deformed Face

# Generator

## Shoulder

- |
- | Upper arm (humerus)

- |
- o Elbow

- |
- | Forearm (radius/ulna)

- |
- o Carpus (wrist) ← bends forward

- |
- | Pastern

- |
- 🐾 Paw

## Hip

- |
- | Femur (thigh)

- |
- o Stifle (knee) ← bends forward

- |
- | Tibia (shin)

- |
- o Hock (ankle) ← bends backward, elevated

- |
- | Metatarsus (rear pastern)

- |
- 🐾 Paw

# Anatomy Map

# Generator

- The negative prompt modifies the noise prediction:
- $\epsilon_{guided} = \epsilon_{uncond} + s \cdot (\epsilon_{cond} - \epsilon_{uncond}) - s_{neg} \cdot (\epsilon_{neg} - \epsilon_{uncond})$
- **Without negative prompt:**  
 $\epsilon = \epsilon_{uncond} + s \cdot (\epsilon_{pos} - \epsilon_{uncond})$
- **With negative prompt:**  
 $\epsilon = \epsilon_{uncond} + s \cdot (\epsilon_{pos} - \epsilon_{uncond}) - s_{neg} \cdot (\epsilon_{neg} - \epsilon_{uncond})$
- Where:
  - $\epsilon_{pos}$  = prediction guided by "golden retriever puppy"
  - $\epsilon_{neg}$  = prediction guided by "cartoon, extra legs, blurry..."
  - $s_{neg}$  = negative guidance scale (usually same as  $s$ )
- **Effect:** Pushes the generation away from the negative concepts while pulling toward positive ones.

# Generator

## Photoreal Post-Processing Pipeline

### **Microfur Luminance Enhancement**

Operates only on the Y-channel of YUV space, enhancing natural fur texture while avoiding color distortion, mimicking real sensor fur granularity.

### **Real Camera Noise Simulation**

Shot-noise and chroma-noise models recreate real sensor imperfections, removing the overly smooth “diffusion look” and increasing perceived realism.

### **Real Tone Curve Application**

Applies DSLR-like S-curve contrast, gentle gamma shift, and dynamic range shaping to approximate professional camera output.

### **Detail Preservation Without Oversharpening**

Adjusts microstructure without edge halos, maintaining real photographic softness while preserving fine detail such as paw edges and muzzle texture.

### **Final Natural-Looking Composition**

The combined post-processing stages align the image with real-world photographic expectations, making the output visually indistinguishable from actual dog photos.

# Generator

## Seed-Controlled Determinism

- **Seed-Controlled Determinism**  
A per-device `torch.Generator` ensures exact reproducibility — same seed, same anatomy, same lighting.
- **Snap-to-8 Resolution Handling**  
Automatically adjusts width & height to multiples of 8 for SDXL compatibility, preventing tensor misalignment.
- **Automated CLI-Based Execution**  
The pipeline supports arguments (`--breed`, `--lora`, `--seed`, `--out`) enabling automated batch generation.
- **Timestamped Image Saving**  
Output file names include timestamps + breed name for dataset consistency and experiment tracking.
- **Production-Ready Modular Design**  
Pipeline is fully modular: anatomy engine, prompt builder, sampler, post-processor, and file saver operate as independent layers for easy extension.

# Findings

- **High-Accuracy Dog Breed Classification Across 120 Classes**  
Achieved **~0.85 accuracy and F1-macro**, demonstrating strong generalization even on visually similar breeds.
- **Realistic, Anatomy-Aware Image Synthesis for 120 Breeds**  
The SDXL + LoRA generator produces coherent, breed-faithful images across multiple visual styles with controlled anatomy and pose.
- **Robust Performance Validated with Multiple Metrics**  
Consistent macro/micro/weighted F1 scores confirm balanced predictions across both rare and common breeds.
- **Unified Recognition–Generation Pipeline**  
Successfully combined ResNet-50 classification with SDXL image generation into a single automated system for end-to-end breed understanding.
- **Interactive Deployment via Streamlit Application**  
Built a full-stack interface enabling real-time breed prediction and customizable dog-image generation for user exploration.

# Conclusion

- **Unified Discriminative + Generative System**

This project successfully integrates a **ResNet-50 breed classifier** with an **SDXL + LoRA image generator**, demonstrating how recognition and synthesis can operate together in a single interactive workflow.

- **Strong Classification Performance**

The classifier achieves **~0.85 macro, micro, and weighted F1**, showing balanced performance across both frequent and rare breeds. Stable micro-accuracy (0.8508) confirms consistent prediction quality on the full 120-breed dataset.

- **Multi-Style, Anatomy-Aware Image Generation**

The SDXL pipeline produces coherent dog images across **five distinct styles**, with LoRA providing controllable stylization. Despite diffusion natural structural challenges, the system maintains breed identity and visual consistency.

- **Identified Limitations**

Current constraints include:

- Classifier restricted to **closed-set 120 breeds**
- Occasional SDXL **anatomical inconsistencies** in limbs and pose
- High GPU memory and compute cost during generation

These highlight clear opportunities for refinement.

- **Future Impact and Extensions**

The project lays a strong foundation for advancements such as **open-set breed classification**, improved **anatomical constraint modules** in diffusion models, cross-breed interpolation, and optimized deployment for real-time applications.

**Link(R): <http://98.83.136.175:8888/>**

Streamlit

**Link: <http://54.162.224.22:8888/>**

**DEMO**

Thank you

